

CS471: Introduction to Artificial Intelligence


Assignment 2: Hill-climbing

Malcolm Zartman


09/12/2025

Question 1: Hill-climbing

- a. Consider a function $f(x) = 2 - x^2$ in the following discrete state-space, where $x \in [-5, 5]$, step-size 0.5. Implement the hill-climbing algorithm in python to find the maximum value for the above function.



- b. Change the step-size to 0.01. Run the hill-climbing algorithm and share your observations.



Colab link:

https://colab.research.google.com/drive/1DoxKoEw_1e0lsEjtlmj_z7lhc-oWU3Eb?usp=sharing

1a)

I first am defining the function for later calculation usage:

```
def f(x):  
    # This is the function we want to maximize:  $f(x) = 2 - x^2$   
    return 2 - (x * x)
```

Below is where I defined my hill_climbing function. It takes one input which is a list of x inputs or state spaces for the f(x) function that we previously defined. In general it follows this algorithm:

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum
 current ← *problem*.INITIAL
 while true do
 neighbor ← a highest-valued successor state of *current*
 if VALUE(*neighbor*) ≤ VALUE(*current*) **then return** *current*
 current ← *neighbor*

For this function, I am however going to be checking both the left and right neighbors to see which neighbor has a greater value, but I am only checking this on the first pass. I am using the variable goingLeft for the direction of the climbing. True if we are going left and decreasing the current_ind by 1, and False if we are going right and increasing the current_ind by 1.

We are starting at the 0th index for no rhyme or reason, but this does allow us to start at the lowest end of the discrete space and allows us to climb to the top over multiple iterations.

Other than tracking and increasing or decreasing the current_ind, depending on the direction of our climbing, we will also be tracking the current_value and at any point that we are not climbing anymore or we are

not finding neighbors that are greater than this current_value, then we will return this current_value as our reasonable solution or local maximum.

```
def hill_climbing(problem_space):
    # Start at the beginning of the problem space
    current_ind = 0 # index of the current state
    current_x = problem_space[current_ind] # first x-value in the
state-space
    current_value = f(current_x)
    # should be -5, -23
    print(f"Starting at x = {current_x}, with f(x) = {current_value}")

    firstPass = True
    goingLeft = False
    neighbor1, neighbor2 = None, None
    while True:
        # for the first pass check both of the neighbors and only use the
greatest one
        if firstPass:
            neighborL = f(problem_space[current_ind-1]) # calculated left
neighbor's value
            neighborR = f(problem_space[current_ind+1]) # calculated right
neighbor's value
            # checking if values were calculated, if so continue - otherwise
return none
            if (neighborL != None) & (neighborR != None):
                if neighborL > neighborR: # if the left neighbor is greater then
we are going left
                    print(f"Going with the Left neighbor at x =
{problem_space[current_ind-1]}, with f(x) = {neighborL}")
                    print(f"Since the Right neighbor at x =
{problem_space[current_ind+1]}, with f(x) = {neighborR} is less than the
Left neighbor")
                    neighbor = neighborL
                    current_ind = current_ind - 1
                    goingLeft = True
                else: # if the right neighbor is greater then we are going right
                    print(f"Going with the Right neighbor at x =
{problem_space[current_ind+1]}, with f(x) = {neighborR}")
```

```

        print(f"Since the Left neighbor at x =
{problem_space[current_ind-1]}, with f(x) = {neighborL} is less than the
Right neighbor")
        neighbor = neighborR
        current_ind = current_ind + 1
        # goingLeft = False # already set
    else:
        print("Calculation Error")
        return None
    firstPass = False
    else: # if it's not the first pass, keep going in the chosen direction
until loc1 max is reached
    if goingLeft:
        neighbor = f(problem_space[current_ind-1])
        current_ind = current_ind - 1
    else:
        neighbor = f(problem_space[current_ind+1])
        current_ind = current_ind + 1

    if neighbor <= current_value: # if the neighbor is not greater, stop
traversing, we found a reasonable solution
        print(f"The neighbor at x = {problem_space[current_ind]}, with f(x)
= {neighbor}, is less than or equal to the current value.")
        print(f"We have reached a reasonable solution at x = {current_x},
with f(x) = {current_value}")
        return current_value
    else: # if the neighbor is greater than the current_value, then keep
searching for a greater one
        print(f"The neighbor at x = {problem_space[current_ind]}, with f(x)
= {neighbor}, is greater than the current value.")
        print("Evaluating next neighbor...")
        current_value = neighbor
        current_x = problem_space[current_ind]

```

In the below code, we are defining the state space or essentially the x-range of our problem, and our step size (the size of our steps in that x-range - determines how many or how little values we are checking as well as how accurate or how long we take in the algorithm). We also are using

the x-range and step-size to create the state space and then send it to the hill_climbing function to find a local_maximum / reasonable solution.

```
# define the discrete state-space for x from -5 to 5 with a step-size of 0.5
x_min = -5.0
x_max = 5.0
step = 0.5

# generate the list of possible x values
states = [x_min + i * step for i in range(int((x_max - x_min) / step) + 1)]

print("Problem: Find the maximum value of f(x) = 2 - x^2")
print(f"State space for x: {states}\n")

# run the hill-climbing algorithm
max_value = hill_climbing(states)

print(f"Maximum function value = {max_value}")
```

Output:

```
➤ Problem: Find the maximum value of f(x) = 2 - x^2
State space for x: [-5.0, -4.5, -4.0, -3.5, -3.0, -2.5, -2.0, -1.5, -1.0, -0.5, 0.0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5.0]

Starting at x = -5.0, with f(x) = -23.0
Going with the Right neighbor at x = -4.5, with f(x) = -18.25
Since the Left neighbor at x = 5.0, with f(x) = -23.0 is less than the Right neighbor
The neighbor at x = -4.5, with f(x) = -18.25, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -4.0, with f(x) = -14.0, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -3.5, with f(x) = -10.25, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -3.0, with f(x) = -7.0, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -2.5, with f(x) = -4.25, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -2.0, with f(x) = -2.0, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -1.5, with f(x) = -0.25, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -1.0, with f(x) = 1.0, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = -0.5, with f(x) = 1.75, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = 0.0, with f(x) = 2.0, is greater than the current value.
Evaluating next neighbor...
The neighbor at x = 0.5, with f(x) = 1.75, is less than or equal to the current value.
We have reached a reasonable solution at x = 0.0, with f(x) = 2.0
Maximum function value = 2.0
```

The maximum value found was 2.0

1b) By changing the step-size to 0.01, it still reached the max value of the $f(x)$ function, but it took a whole lot longer to reach that solution. So much so, that it wouldn't make sense to paste the output of it into this document.

```
# 1b) STEP SIZE OF 0.01

# define the discrete state-space for x from -5 to 5 with a step-size of
0.5
x_minb = -5.0
x_maxb = 5.0
stepb = 0.01

# generate the list of possible x values
statesb = [x_minb + i * stepb for i in range(int((x_maxb - x_minb) /
stepb) + 1)]

print("Problem: Find the maximum value of  $f(x) = 2 - x^2$ ")
print(f"State space for x: {statesb}\n")


# run the hill-climbing algorithm
max_valueb = hill_climbing(statesb)

print(f"Maximum function value = {max_valueb}")
```

Question 2: Random-restart hill-climbing

a. Consider a function

$g(x) = (0.0051x^5) - (0.1367x^4) + (1.24x^3) - (4.456x^2) + (5.66x) - 0.287$ in the following discrete state-space, where $x \in [0, 10]$, step-size 0.5. Implement the random-restart hill-climbing algorithm for 20 random restarts in python to find the global maximum value for the above function.



- b. Run the hill-climbing algorithm for the function $g(x)$. Compare and analyze the results of hill-climbing with the random-restart hill-climbing algorithm.
-

Colab link:

https://colab.research.google.com/drive/1DoxKoEw_1e0lsEjtlmj_z7lhc-oWU3Eb?usp=sharing

2a)

First I am defining the function for later calculation usages:

```
def g(x):
    # This is the function we want to maximize for the random_restart
    algorithm:  $f(x) = 2 - x^2$ 
    return (0.0051 * x**5) - (0.1367 * x**4) + (1.24 * x**3) - (4.456 *
x**2) + (5.66 * x) - 0.287
```

Below is the code for the random_restart function or the random-restart hill climbing algorithm. This differs from the previous algorithm because instead of starting at a predefined location, we are starting at a random location and then running the algorithm a certain amount of times. In my “main” code I define the looping logic for the iteration portion of the algorithm, but in the function below, the major difference in terms of code from the previous hill climbing function, is the random location generation and the

handling for it. I use the variable `space_length` for generating the random locations, and to also ensure I don't have an index that is out of bounds with my `problem_space` list by using the modulus in order to loop back around when iterating through the neighbors.

```
import random
def random_restart(problem_space):
    # ONLY CHANGE: Start at a random state in the problem space
    space_length = len(problem_space) # will use space_length & modulus to
    ensure no 'index out of range' errors
    current_ind = random.randint(0,space_length-1) # index of the current
    state
    current_x = problem_space[current_ind] # first x-value in the
    state-space
    current_value = g(current_x)
    print(f"Starting at x = {current_x}, with f(x) = {current_value}")

    firstPass = True
    goingLeft = False
    neighborL, neighborR = None, None
    while True:
        # for the first pass check both of the neighbors and only use the
        greatest one
        if firstPass:
            neighborL =
            f(problem_space[(current_ind-1+space_length)%space_length]) # calculated
            left neighbor's value
            neighborR = f(problem_space[(current_ind+1)%space_length]) #
            calculated right neighbor's value
            # checking if values were calculated, if so continue - otherwise
            return none
            if (neighborL != None) & (neighborR != None):
                if neighborL > neighborR: # if the left neighbor is greater then
                we are going left
                    print(f"Going with the Left neighbor at x =
                    {problem_space[(current_ind-1+space_length)%space_length]}, with f(x) =
                    {neighborL}")
                    print(f"Since the Right neighbor at x =
                    {problem_space[(current_ind+1)%space_length]}, with f(x) = {neighborR} is
                    less than the Left neighbor")
                    neighbor = neighborL
```



```

        current_ind = current_ind - 1
        goingLeft = True
    else: # if the right neighbor is greater then we are going right
        print(f"Going with the Right neighbor at x = {problem_space[(current_ind+1)%space_length]}, with f(x) = {neighborR}")
        print(f"Since the Left neighbor at x = {problem_space[(current_ind-1+space_length)%space_length]}, with f(x) = {neighborL} is less than the Right neighbor")
        neighbor = neighborR
        current_ind = current_ind + 1
        # goingLeft = False # already set
    else:
        print("Calculation Error")
        return None
    firstPass = False
    else: # if it's not the first pass, keep going in the chosen direction
        until loc1 max is reached
        if goingLeft:
            neighbor = f(problem_space[(current_ind-1+space_length)%space_length])
            current_ind = current_ind - 1
        else:
            neighbor = f(problem_space[(current_ind+1)%space_length])
            current_ind = current_ind + 1

    if neighbor <= current_value: # if the neighbor is not greater, stop
        traversing, we found a reasonable solution
        print(f"The neighbor at x = {problem_space[current_ind%space_length]}, with f(x) = {neighbor}, is less
        than or equal to the current value.")
        print(f"We have reached a reasonable solution at x = {current_x},
        with f(x) = {current_value}")
        return current_value
    else: # if the neighbor is greater than the current_value, then keep
        searching for a greater one
        print(f"The neighbor at x = {problem_space[current_ind%space_length]}, with f(x) = {neighbor}, is
        greater than the current value.")
        print("Evaluating next neighbor...")
        current_value = neighbor

```

```
current_x = problem_space[current_ind%space_length]
```

Below is the logic for defining the state space as well as for the rest of the random restart hill-climbing algorithm. Here is where I handle the amount of iterations of calling this algorithm and then finding an ultimate “local-maximum” or “reasonable solution.” The desired iteration amount was 20, and as I call the algorithm, I save the current local_max and then check if it was greater than the previously found max, if so, I save it as the new max in max_found, otherwise I keep looping. That process continues until I call the algorithm 20 times. At the end I display the max_found.

```
# define the discrete state-space for x from 0 to 10 with a step-size of
0.5
x_min2 = 0.0
x_max2 = 10.0
step2 = 0.5

# generate the list of possible x values
states2 = [x_min2 + i * step2 for i in range(int((x_max2 - x_min2) /
step2) + 1)]

print("Problem: Find the maximum value of  $g(x) = (0.0051x^5) - (0.1367x^4) +$ 
 $(1.24x^3) - (4.456x^2) + (5.66x) - 0.287$ ")
print(f"State space for x: {states2}\n")

# run the random restart hill-climbing algorithm
local_max = None
max_found = None
iterations = 20
count = 0
while count < iterations:
    count = count + 1
    print(f"Iteration: {count}")
    max_value2 = random_restart(states2)
    local_max = max_value2
    if (max_found == None) or (local_max > max_found):
        max_found = local_max
    print("\n")

print(f"\nMaximum function value (rounded by 4) = {round(max_found,4)}")
```

Output:

Problem: Find the maximum value of $g(x) = (0.0051x^5) - (0.1367x^4) + (1.24x^3) - (4.456x^2) + (5.66x) - 0.287$

State space for x: [0.0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5.0, 5.5, 6.0, 6.5, 7.0, 7.5, 8.0, 8.5, 9.0, 9.5, 10.0]

Iteration: 1

Starting at x = 5.0, with $f(x) = 2.1130000000000093$

Going with the Left neighbor at x = 4.5, with $f(x) = -18.25$

Since the Right neighbor at x = 5.5, with $f(x) = -28.25$ is less than the Left neighbor

The neighbor at x = 4.5, with $f(x) = -18.25$, is less than or equal to the current value.

We have reached a reasonable solution at x = 5.0, with $f(x) = 2.1130000000000093$

Iteration: 2

Starting at x = 2.0, with $f(x) = 1.1049999999999995$

Going with the Left neighbor at x = 1.5, with $f(x) = -0.25$

Since the Right neighbor at x = 2.5, with $f(x) = -4.25$ is less than the Left neighbor

The neighbor at x = 1.5, with $f(x) = -0.25$, is less than or equal to the current value.

We have reached a reasonable solution at x = 2.0, with $f(x) = 1.1049999999999995$

Iteration: 3

Starting at x = 7.5, with $f(x) = 3.1360468750000257$

Going with the Left neighbor at x = 7.0, with $f(x) = -47.0$

Since the Right neighbor at x = 8.0, with $f(x) = -62.0$ is less than the Left neighbor

The neighbor at x = 7.0, with $f(x) = -47.0$, is less than or equal to the current value.

We have reached a reasonable solution at x = 7.5, with $f(x) = 3.1360468750000257$

Iteration: 4

Starting at x = 5.0, with $f(x) = 2.1130000000000093$

Going with the Left neighbor at x = 4.5, with $f(x) = -18.25$

Since the Right neighbor at x = 5.5, with $f(x) = -28.25$ is less than the Left neighbor

The neighbor at x = 4.5, with $f(x) = -18.25$, is less than or equal to the current value.

We have reached a reasonable solution at x = 5.0, with $f(x) = 2.1130000000000093$

Iteration: 5

Starting at $x = 0.0$, with $f(x) = -0.287$

Going with the Right neighbor at $x = 0.5$, with $f(x) = 1.75$

Since the Left neighbor at $x = 10.0$, with $f(x) = -98.0$ is less than the Right neighbor

The neighbor at $x = 0.5$, with $f(x) = 1.75$, is greater than the current value.

Evaluating next neighbor...

The neighbor at $x = 1.0$, with $f(x) = 1.0$, is less than or equal to the current value.

We have reached a reasonable solution at $x = 0.5$, with $f(x) = 1.75$

Iteration: 6

Starting at $x = 5.0$, with $f(x) = 2.1130000000000093$

Going with the Left neighbor at $x = 4.5$, with $f(x) = -18.25$

Since the Right neighbor at $x = 5.5$, with $f(x) = -28.25$ is less than the Left neighbor

The neighbor at $x = 4.5$, with $f(x) = -18.25$, is less than or equal to the current value.

We have reached a reasonable solution at $x = 5.0$, with $f(x) = 2.1130000000000093$

Iteration: 7

Starting at $x = 6.0$, with $f(x) = 3.59139999999999637$

Going with the Left neighbor at $x = 5.5$, with $f(x) = -28.25$

Since the Right neighbor at $x = 6.5$, with $f(x) = -40.25$ is less than the Left neighbor

The neighbor at $x = 5.5$, with $f(x) = -28.25$, is less than or equal to the current value.

We have reached a reasonable solution at $x = 6.0$, with $f(x) = 3.59139999999999637$

Iteration: 8

Starting at $x = 0.0$, with $f(x) = -0.287$

Going with the Right neighbor at $x = 0.5$, with $f(x) = 1.75$

Since the Left neighbor at $x = 10.0$, with $f(x) = -98.0$ is less than the Right neighbor

The neighbor at $x = 0.5$, with $f(x) = 1.75$, is greater than the current value.

Evaluating next neighbor...

The neighbor at $x = 1.0$, with $f(x) = 1.0$, is less than or equal to the current value.

We have reached a reasonable solution at $x = 0.5$, with $f(x) = 1.75$

Iteration: 9

Starting at $x = 7.0$, with $f(x) = 3.8080000000000056$

Going with the Left neighbor at $x = 6.5$, with $f(x) = -40.25$
Since the Right neighbor at $x = 7.5$, with $f(x) = -54.25$ is less than the Left neighbor
The neighbor at $x = 6.5$, with $f(x) = -40.25$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 7.0$, with $f(x) = 3.808000000000056$

Iteration: 10

Starting at $x = 3.0$, with $f(x) = 0.23559999999999354$
Going with the Left neighbor at $x = 2.5$, with $f(x) = -4.25$
Since the Right neighbor at $x = 3.5$, with $f(x) = -10.25$ is less than the Left neighbor
The neighbor at $x = 2.5$, with $f(x) = -4.25$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 3.0$, with $f(x) = 0.23559999999999354$

Iteration: 11

Starting at $x = 6.0$, with $f(x) = 3.59139999999999637$
Going with the Left neighbor at $x = 5.5$, with $f(x) = -28.25$
Since the Right neighbor at $x = 6.5$, with $f(x) = -40.25$ is less than the Left neighbor
The neighbor at $x = 5.5$, with $f(x) = -28.25$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 6.0$, with $f(x) = 3.59139999999999637$

Iteration: 12

Starting at $x = 1.5$, with $f(x) = 1.708684374999998$
Going with the Left neighbor at $x = 1.0$, with $f(x) = 1.0$
Since the Right neighbor at $x = 2.0$, with $f(x) = -2.0$ is less than the Left neighbor
The neighbor at $x = 1.0$, with $f(x) = 1.0$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 1.5$, with $f(x) = 1.708684374999998$

Iteration: 13

Starting at $x = 1.5$, with $f(x) = 1.708684374999998$
Going with the Left neighbor at $x = 1.0$, with $f(x) = 1.0$
Since the Right neighbor at $x = 2.0$, with $f(x) = -2.0$ is less than the Left neighbor
The neighbor at $x = 1.0$, with $f(x) = 1.0$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 1.5$, with $f(x) = 1.708684374999998$

Iteration: 14

Starting at $x = 9.0$, with $f(x) = -2.06179999999999705$

Going with the Left neighbor at $x = 8.5$, with $f(x) = -70.25$
Since the Right neighbor at $x = 9.5$, with $f(x) = -88.25$ is less than the Left neighbor
The neighbor at $x = 8.5$, with $f(x) = -70.25$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 9.0$, with $f(x) = -2.0617999999999705$

Iteration: 15

Starting at $x = 7.5$, with $f(x) = 3.1360468750000257$
Going with the Left neighbor at $x = 7.0$, with $f(x) = -47.0$
Since the Right neighbor at $x = 8.0$, with $f(x) = -62.0$ is less than the Left neighbor
The neighbor at $x = 7.0$, with $f(x) = -47.0$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 7.5$, with $f(x) = 3.1360468750000257$

Iteration: 16

Starting at $x = 9.5$, with $f(x) = -4.3277656249999845$
Going with the Left neighbor at $x = 9.0$, with $f(x) = -79.0$
Since the Right neighbor at $x = 10.0$, with $f(x) = -98.0$ is less than the Left neighbor
The neighbor at $x = 9.0$, with $f(x) = -79.0$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 9.5$, with $f(x) = -4.3277656249999845$

Iteration: 17

Starting at $x = 6.5$, with $f(x) = 3.9287781250000213$
Going with the Left neighbor at $x = 6.0$, with $f(x) = -34.0$
Since the Right neighbor at $x = 7.0$, with $f(x) = -47.0$ is less than the Left neighbor
The neighbor at $x = 6.0$, with $f(x) = -34.0$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 6.5$, with $f(x) = 3.9287781250000213$

Iteration: 18

Starting at $x = 1.5$, with $f(x) = 1.708684374999998$
Going with the Left neighbor at $x = 1.0$, with $f(x) = 1.0$
Since the Right neighbor at $x = 2.0$, with $f(x) = -2.0$ is less than the Left neighbor
The neighbor at $x = 1.0$, with $f(x) = 1.0$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 1.5$, with $f(x) = 1.708684374999998$

Iteration: 19

Starting at $x = 3.0$, with $f(x) = 0.23559999999999354$

Going with the Left neighbor at $x = 2.5$, with $f(x) = -4.25$
Since the Right neighbor at $x = 3.5$, with $f(x) = -10.25$ is less than the Left neighbor
The neighbor at $x = 2.5$, with $f(x) = -4.25$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 3.0$, with $f(x) = 0.235599999999999354$

Iteration: 20

Starting at $x = 8.0$, with $f(x) = 1.88260000000000311$
Going with the Left neighbor at $x = 7.5$, with $f(x) = -54.25$
Since the Right neighbor at $x = 8.5$, with $f(x) = -70.25$ is less than the Left neighbor
The neighbor at $x = 7.5$, with $f(x) = -54.25$, is less than or equal to the current value.
We have reached a reasonable solution at $x = 8.0$, with $f(x) = 1.88260000000000311$

Maximum function value (rounded by 4) = 3.9288

2b) By using the hill-climbing algorithm for the $g(x)$ function, it evaluated the function rather fast, however, the result was not the global maximum of the function, rather it was a local maximum which is still a reasonable solution, just not the one we are ultimately looking for. So the random-restart hill climbing algorithm is for sure a better solution for a more complex function such as $g(x)$, in which we run it multiple times and only keep the ultimate maximum out of the ones we find. And in this scenario, the random-restart was able to find the global maxim while the regular hill-climbing function only found the local-maximum. With this algorithm a way to optimize it further is by figuring out what would be a good amount of iterations to run that random-restart algorithm for.

```
# 2b) HILL CLIMBING WITH G(x)

# define the discrete state-space for x from -5 to 5 with a step-size of
0.5
x_min2b = -5.0
x_max2b = 5.0
step2b = 0.5

# generate the list of possible x values
states2b = [x_min2b + i * step2b for i in range(int((x_max2b - x_min2b) /
step2b) + 1)]
```

```

print("Problem: Find the maximum value of  $f(x) = 2 - x^2$ ")
print(f"State space for x: {states2b}\n")

# run the hill-climbing algorithm
max_value2b = hill_climbingG(states2b)

print(f"Maximum function value = {max_value2b}")

```

Output:

```

➤ Problem: Find the maximum value of  $f(x) = 2 - x^2$ 
State space for x: [-5.0, -4.5, -4.0, -3.5, -3.0, -2.5, -2.0, -1.5, -1.0, -0.5, 0.0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5.0]

Starting at x = -5.0, with  $g(x) = -396.36199999999997$ 
Going with the Left neighbor at x = 5.0, with  $g(x) = 2.1130000000000093$ 
Since the Right neighbor at x = -4.5, with  $g(x) = -294.452478125$  is less than the Left neighbor
The neighbor at x = 5.0, with  $g(x) = 2.1130000000000093$ , is greater than the current value.
Evaluating next neighbor...
The neighbor at x = 4.5, with  $g(x) = 1.2993906249999993$ , is less than or equal to the current value.
We have reached a reasonable solution at x = 5.0, with  $g(x) = 2.1130000000000093$ 
Maximum function value = 2.1130000000000093

```